

2. Evolutionary computing and the simple genetic algorithm

What is evolutionary computing?

- Evolutionary Computing (EC) techniques:
 - are a diverse collection of 'nature-inspired' techniques
 - heuristic: not guaranteed to find the best solution
 - the opposite of an 'exact' algorithm
 - stochastic: they use random number generators
 - the opposite of deterministic algorithms
 - intended for hard optimization problems (e.g., NP-hard)
- Some examples:
 - Genetic Algorithms (GA)
 - Genetic Programming (GP)
 - Evolutionary Strategies (ES)
 - ...

What is evolutionary computing?

- In contrast, many other heuristic algorithms do not use a population of solutions; e.g.,
 - Hillclimbing
 - Simulated annealing
 - Tabu search
- ...

Simple Hillclimbing

New current solution = the first fitter point it finds in the neighbourhood of the current point

Start with any point p

repeat

$q = \text{chooseNextPoint}(p)$

 if $p == q$ then return p as final solution


$\text{chooseNextPoint}(p)$:

for each neighbour q of p

 if q is fitter than p return q

return p

Simple Hillclimbing



Steepest ascent hillclimbing

Evaluates all neighbours of the current point and selects the fittest of them

```
start with any point p
```

```
repeat
```

```
    q = chooseNextPoint(p)
```

```
    if p == q then return p as final solution
```

```
chooseNextPoint(p):
```


```
    fittest = p
```

```
    for each neighbour q of p
```

```
        if q is fitter than fittest then fittest = q
```

```
    return fittest
```

Steepest ascent hillclimbing



Simulated annealing

Randomly choose a different solution that is better than current solution, or even worse than current solution if temperature is higher.

Gradually reduce temperature.

https://en.wikipedia.org/wiki/Simulated_annealing

The explore/exploit balance

Hillclimbing **exploits** current knowledge of search space.

Simulated annealing also **explores** search space.

The explore/exploit balance is very important in RL (later).

Disadvantages of EC

- EC is heuristic
 - not guaranteed to find best possible solution
- EC is stochastic
 - can produce a different solution every time we run it
 - we often want to run it repeatedly in case the next run returns a better solution
 - how many times should we run it?
- Evolutionary algorithms are complex systems
 - complex or chaotic behaviour can emerge from simple interactions between simple parts
 - very hard to predict how overall system will behave
 - one approach is to run algorithm many times and see what happens
 - but that can be slow, especially if you can change many parameters
- Given these disadvantages, why use them?

Computational compromises

- Instead of EC, we can use a simple algorithm which is guaranteed to find the best possible solution:
 bestSolution = null;
 for *each solution s in possibleSolutions* do
 if isBetter(s, bestSolution) then bestSolution = s;
 end
- If this is fast enough for your problem, use it!
 - But for many problems, it scales too slowly
- For many computer science problems there is a simple, exact algorithm ... which is too slow to use
- So instead we either:
 1. Accept an approximate (suboptimal) solution
 2. Or find an exact solution to an approximation of the original problem
- EC takes the first approach

Alternatives to EC

- Many other optimization methods exist
 - Some are very efficient
 - Some are also exact, or at least guaranteed to improve the longer you run them
 - But they may require information about the problem that you do not have
 - We'll return to this when we cover Dynamic Programming
 - Or require conditions that do not hold
 - E.g., that a function be differentiable

Alternatives to EC

- Nonetheless, there are *many* useful optimization methods, which we will not cover; e.g.,
 - deterministic search algorithms (e.g., depth-first search)
 - linear programming
 - integer programming
 - constraint satisfaction
 - optimal control
 - ... and many more. See mathematical optimization in wikipedia for an overview
- Always use the right tool for the job

Darwin, Mendel, and the modern synthesis

EC is inspired by 'neo-Darwinian' evolutionary theory

- 1859: publication of *Origin of Species* by Charles Darwin
 - Explained evolution as due to natural selection acting on heritable variation
 - Very similar ideas proposed by A. R. Wallace at about the same time
- 1865 - Gregor Mendel presented his work on *Experiments in Plant Hybridisation*
 - He demonstrated the particulate nature of inheritance (i.e., genes are discrete particles)
 - Darwin had assumed genes were blended between parents (like averaging two numbers)
- These two achievements together paved the way for the 'Modern Synthesis'...

The modern synthesis (neo-Darwinism)

- One problem that had worried Darwin was 'regression' of traits
 - Darwin assumed blending of heritable material
 - Hence the 'value' of any trait would tend to converge in a population as evolution progressed
 - The resulting lack of variation in the population would give natural selection nothing to act on, so evolution would stall
- Mendel's results went unnoticed by evolutionists for about 30 years

The modern synthesis (neo-Darwinism)

- Once rediscovered, Mendel's laws of particulate (i.e., genetic) inheritance were incorporated into Darwinian evolutionary theory
- The result was the 'modern synthetic theory of evolution', or 'neo-Darwinism'
 - Natural selection acting on genetic variation in populations
 - Primarily mathematical in nature, modelling gene spread in populations
 - Population genetics
 - Quantitative genetics
- Subsequently, the physical basis of particulate inheritance was found
 - Discovery of DNA by James Watson and Francis Crick, 1962
Nobel Laureates in Medicine

The emergence of evolutionary computing

- Evolutionary ideas were being applied in optimization as early as the mid 20th century; e.g.,
 - Box (1957) Evolutionary operation: a method for increasing industrial productivity. Applied Statistics 6, 81-101
 - Bremermann (1962) Optimization through evolution and recombination. In: Self-Organizing Systems.
- 1970s: Genetic Algorithms became a prominent research area
 - John Holland presented a mathematical definition of GAs, and theory explaining their performance
 - Holland (1975) Adaptation in Natural and Artificial Systems
 - Holland was mainly interested in adaptive systems, but his student Ken De Jong catalysed research on GAs for optimization, formulating a test suite of problems
 - De Jong (1975) An Analysis of the Behavior of a Class of Genetic Adaptive Systems. PhD thesis, University of Michigan

The simple genetic algorithm

- What is the 'Simple Genetic Algorithm' ?
- Genetic algorithms have many variations on a central theme
 - Choice of selection operator
 - Choice of genetic operators
 - ...
- The SGA is the canonical description of what a GA is
 - Here we will describe the SGA according to Holland's original proposals
- As an aside, the term “SGA” also refers to Michael Vose's mathematical description of a GA
 - Vose (1999) The Simple Genetic Algorithm: Foundation and Theory

Terminology

- *Chromosome* - a solution to the problem the GA is solving, encoded as a fixed length string in some finite alphabet
- *Gene* - a position on the chromosome which can take a particular value
- *Locus* - see gene
- *Allele* - a value that a gene can take (e.g., 0 or 1 in binary chromosomes)
- *Population* - the set of chromosomes the GA is acting on
- *Fitness* - the objective value of the solution a chromosome encodes for a problem
- *Fitness/objective function* - the function that returns the objective value of a given solution

Terminology

- *Selection* - the mechanism for choosing chromosomes to reproduce, based on their fitness
- *Parent* - a chromosome selected for reproduction
- *Offspring* - the chromosome(s) resulting from reproduction of one or more parents
- *Crossover* - the process of recombining alleles from multiple parents during reproduction to produce offspring
- *Mutation* - the process of randomly replacing offsprings' alleles with randomly chosen alternatives during reproduction

The simple GA

- The Simple GA

- generate initial population;

- while *termination criterion unsatisfied* do

- evaluate population;

- select parents;

- reproduce using genetic operators;

- replace population with offspring;

- end

- Genetic operators

- Single-point crossover - recombines alleles from two parents into offspring

- Point mutation - randomly changes alleles in offspring at a very low rate

Selection

- A key idea of the GA is that parents should be selected based on their fitness
- Simplest way of doing this is using roulette wheel selection
 - Each chromosome i in population P is given a probability of selection p_i based on its fitness f_i as fraction of total population fitness:

$$p_i = \frac{f_i}{\sum_{j \in P} f_j}$$

- So for a given population we construct a single-pointer 'roulette wheel' where each chromosome's proportion of the wheel is given by the equation above
- To select a parent for reproduction we uniformly sample from $[0, 1)$
 - i.e., we spin the roulette wheel once and see which chromosome its pointer indicates

Selection

- Having selected a parent, we then reproduce that parent according to our operators, parameters, etc.
 - According to our parameters, we may apply crossover, and/or mutation, and/or any other operators
 - Often crossover rate (probability of applying crossover during reproduction) is denoted by χ
 - Similarly mutation rate (per gene probability of mutation during reproduction) by μ
- Usually population size is held constant...
- ...we continue selecting and reproducing until we have enough offspring to replace the current generation
 - Hence such a GA is called a 'generational' GA

Representation

- The SGA uses binary strings
 - The simplest representation
 - Widely used
 - Have some nice theoretical properties
- Other algorithms also use:
 - higher cardinality alphabets
 - more complex data structures: trees, graphs

Single-point crossover

- Holland's initial proposal for genetic recombination in GAs was single-point crossover, sometimes denoted 1X
- For a pair of parent chromosomes of length l , we (uniformly) choose at random an integer from $\{1, 2, \dots, l-1\}$
- We then exchange portions of each chromosome, combining the substring left of the cutpoint from one chromosome with the substring right of the cutpoint from the other chromosome, and vice versa
 - e.g., the parent chromosomes 11111 and 00000 with the randomly selected crossover point at position 3...
 - ...yield the offspring chromosomes 11100 and 00011
- We now have two offspring chromosomes, each containing alleles from both parent chromosomes
- Optionally we can discard one of the offspring if we only require one offspring

Point mutation

- During reproduction, mutation occurs at a (usually low) rate, in addition to crossover
- Normally mutation rate is expressed as a per-gene probability
 - Called “point mutation” as we mutate 1 point (gene) at a time
- If a gene is mutated, its current allele is exchanged for a different randomly selected allele
 - If our chromosomes are binary strings, there is only one other possible allele and the gene is simply bit-flipped
- Typical mutation rate about 0.01 probability of mutation per gene `transcribed' (copied) from parent to offspring

A simple example - 1-max

- Let's work through a simple example of the SGA solving a problem
 - Problem: 1-Max (binary fixed-length chromosome; fitness is number of ones in chromosome)
 - 1-Max is not a good problem to apply a GA to as the fitness function is linearly separable (e.g., hill-climbing will do much better), but it is often used for illustration
 - Chromosome length (l) = 5
 - Population size (N) = 4
 - Crossover rate (χ) = 1.0
 - Mutation rate (μ) = 0.25 (which is unrealistically high for most problems)